



MongoDB

MapReduce



MapReduce

- But first, some functional programming
- Using Ruby blocks/closures
- Will help understand how MapReduce works



MapReduce



MapReduce

Functional programming for PHP
programmers.



MapReduce

Functional programming for PHP programmers.

Javascripters, Pythonistas and Rubyists feel free to doze off for a bit :P



MapReduce

- What is MapReduce?
- A method of aggregation that can easily be parallelized across multiple servers
 - It splits up a problem
 - Sends chunks of it to different machines
 - Lets each machine solve its part of the problem
- When all machines are finished, they merge all the pieces of the solution back into a full solution



MapReduce

- Caveats:
- Can be slow
- Not meant to be used “real-time”
- Usually run as a background job, which creates a collection of results, which you then query
- Think batch processing



MapReduce

Example: count number of tags

```
db.things.insert( { _id : 1, tags : ['dog',  
'cat'] } );
```

```
db.things.insert( { _id : 2, tags : ['cat'] } );
```

```
db.things.insert( { _id : 3, tags : ['mouse', 'cat',  
'dog'] } );
```

```
db.things.insert( { _id : 4, tags : [] } );
```




MapReduce

```
> mr = db.things.mapReduce(m, r, { out : "myoutput" } );  
{  
  "result" : "myoutput",  
  "timeMillis" : 12,  
  "counts" : {  
    "input" : 4,  
    "emit" : 6,  
    "output" : 3  
  },  
  "ok" : 1,  
}
```



MapReduce

- "result" : "myoutput",
- The name of the collection the MapReduce results were stored in
- "timeMillis" : 12
- How long the operation took, in milliseconds



MapReduce

- "counts" : { ... }
- This embedded document contains three keys:
 - "input" : 6
The number of documents sent to the map function.
 - "emit" : 14
The number of times emit was called in the map function.
 - "output" : 5
- The number of documents created in the result collection. "counts" is mostly useful for debugging.